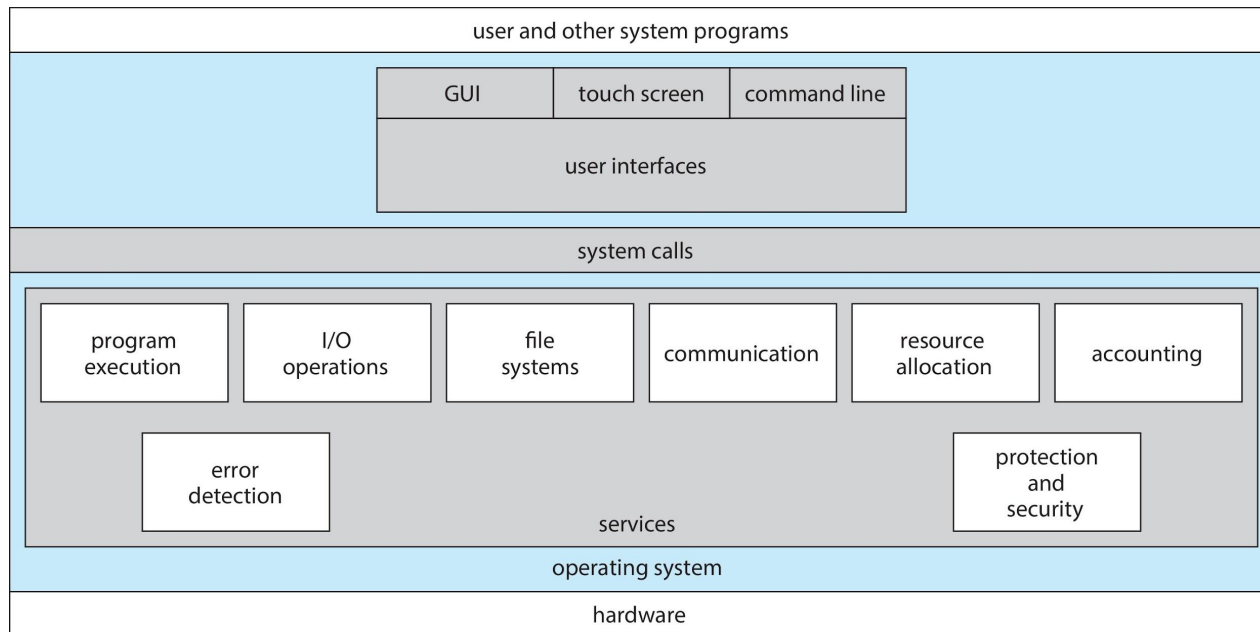



# Review: A View of Operating System Services



# Review: API – System Call – OS Relationship



# Review: System Call Parameter Passing

- Three general methods used to pass parameters to the OS
  - pass the parameters in registers
    - Simplest (no context copy)
    - In some cases, may be more parameters than registers
  - address of a block passed as a parameter in a register
    - Parameters stored in the block, or table, in memory
    - This approach taken by Linux and Solaris
  - Parameters placed, or pushed, onto the stack by the program and popped off the stack by the operating system
- Block and stack methods do not limit the number or length of parameters being passed
  - however they use memory

# Review: kernel implementation structure

	<b>monolithic kernel</b>	<b>microkernel</b>
<b>Reliability</b>	If one driver crashes, entire kernel fails	Kernel can restart system program as needed
<b>Ease of Development</b>	Must get entire kernel to work	Able to test just one part without affecting rest
<b>Speed</b>	faster: No extraneous context switching,	Slow: due to message passing (context switches)
<b>memory</b>	Relatively modest in memory usage	Memory footprint much larger


# hybrid kernel

- In practice, modern OSes are hybrid
- Linux is more monolithic, than current Windows and macOS
  - <https://makelinux.github.io/kernel/diagram/> I
  - [https://en.wikipedia.org/wiki/Architecture\\_of\\_macOS](https://en.wikipedia.org/wiki/Architecture_of_macOS)
  - <https://docs.microsoft.com/en-us/windows-hardware/drivers/kernel/overview-of-windows-components>
- Loadable kernel modules: bit of code that kernel can load (and usually unload) while system is running, to extend functionality
- Examples:
  - Linux kernel modules (lkm),
  - Windows device drivers
  - macOS extension

# Process Management

## Chapter 3: Processes

---



# Outline

- Process Concept
- Process Scheduling
- Operations on Processes
- Interprocess Communication
- IPC in Shared-Memory Systems
- IPC in Message-Passing Systems
- Examples of IPC Systems
- Communication in Client-Server Systems

# Process Concept


- An operating system executes a variety of programs that run as a process.
- **Process** – a program in execution; process execution must progress in sequential fashion. No parallel execution of instructions of a single process
- Multiple parts
  - The program code, also called **text section**
  - Current activity including **program counter**, processor registers
  - **Stack** containing temporary data
    - Function parameters, return addresses, local variables
  - **Data section** containing global variables
  - **Heap** containing memory dynamically allocated during run time



# Process Concept (Cont.)

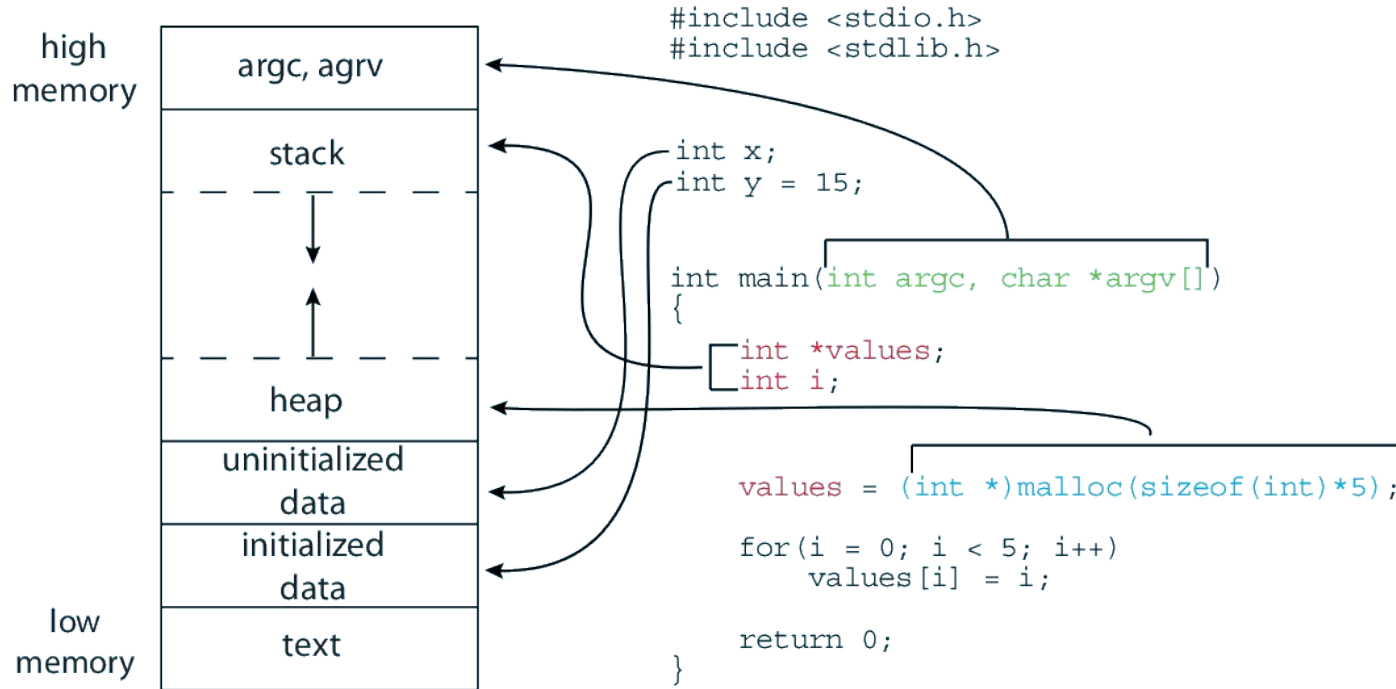
- Program is **passive** entity stored on disk (**executable file**); process is **active**
  - Program becomes process when an executable file is loaded into memory
- Execution of program started via GUI mouse clicks, command line entry of its name, etc.
- One program can be several processes
  - Consider multiple users executing the same program

# Process in Memory (user view of a process)



**Figure 3.1** Layout of a process in memory.

# Memory Layout of a C Program




```
$ gcc memory.c -o memory
```

```
$ size memory
```

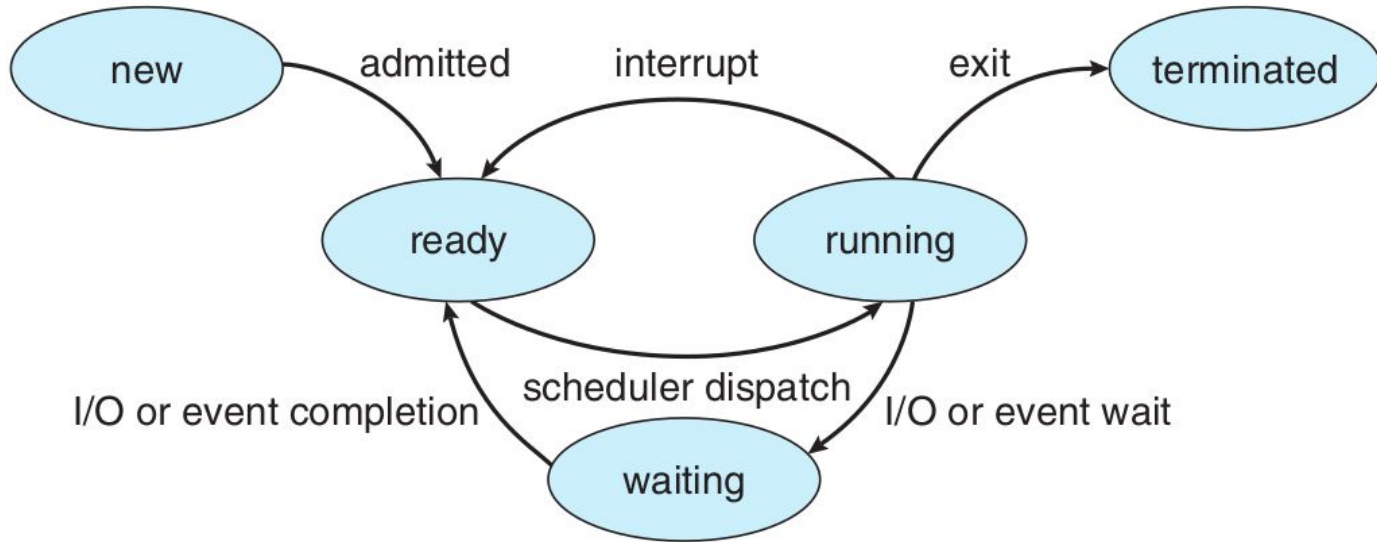
text	data	bss	dec	hex	filename
1603	600	8	2211	8a3	memory

# Kernel view of process

## Process Control Block (PCB)



# Process States



**Figure 3.2** Diagram of process state.


# Process State

- As a process executes, it changes **state**
  - **New:** The process is being created
  - **Running:** Instructions are being executed
  - **Waiting:** The process is waiting for some event to occur
  - **Ready:** The process is waiting to be assigned to a processor
  - **Terminated:** The process has finished execution

# Process Control Block (PCB)

Information associated with each process(also called **task control block**)

- Process state
  - – running, waiting, etc.
- Program counter
  - – location of instruction to next execute
- CPU registers
  - – contents of all process-centric registers
- CPU scheduling information
  - - priorities, scheduling queue pointers
- Memory
  - -management information – memory allocated to the process
- Accounting information
  - – CPU used, clock time elapsed since start, time limits
- I/O status information
  - – I/O devices allocated to process, list of open files



# Threads

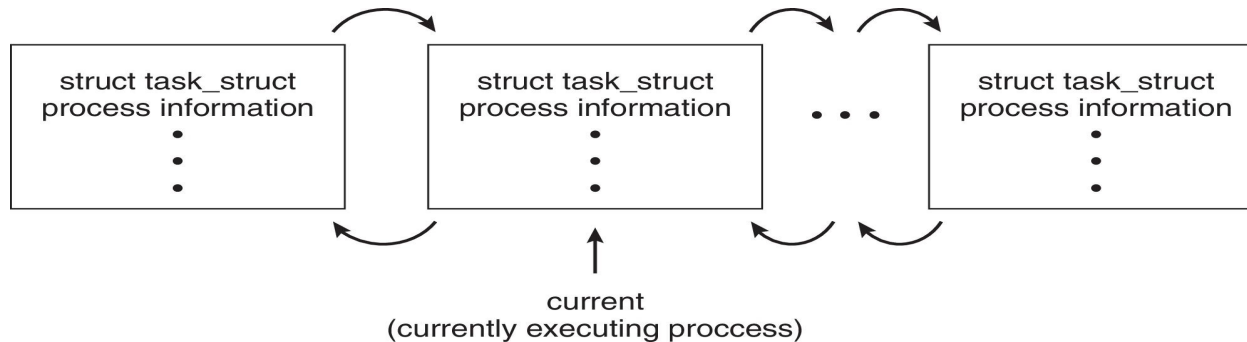
- The process model so far performs **a single thread of execution**.
- Consider having **multiple program counters per process**
  - Multiple locations can execute at once
    - Multiple threads of control -> **threads**
- Must then have storage for thread details, multiple program counters in PCB



# Process Representation in Linux

Represented by the C structure `task_struct`

```
pid t_pid;                /* process identifier */
long state;               /* state of the process */
unsigned int time_slice;  /* scheduling information */
struct task_struct *parent; /* this process's parent */
struct list_head children; /* this process's children */
struct files_struct *files; /* list of open files */
struct mm_struct *mm;     /* address space of this process */
```



# Process Scheduling

**Process scheduler** selects among available processes for next execution on CPU core



- Goal
  - Maximize CPU use (**CPU utilization**),
    - quickly switch processes onto CPU core
  - Switch a CPU core among processes so frequently that users can interact with each program while it is running .

# Process Scheduling-implementation

- Maintains **scheduling queues** of processes
  - **Ready queue** – set of all processes residing in main memory, ready and waiting to execute
  - **Wait queues** – set of processes waiting for an event (i.e., I/O)
  - Processes migrate among the various queues

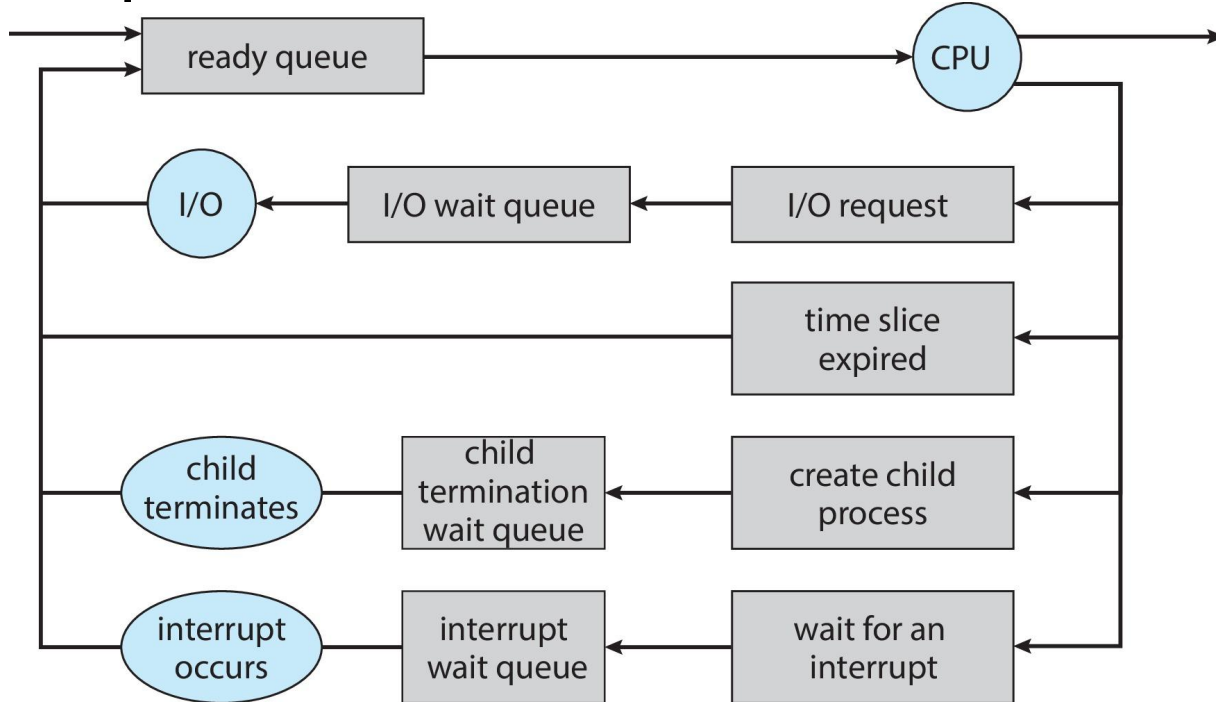
# Ready and Wait Queues

- As processes enter the system, they are put into a **ready queue**,
  - where they are ready and waiting to execute on a CPU's core



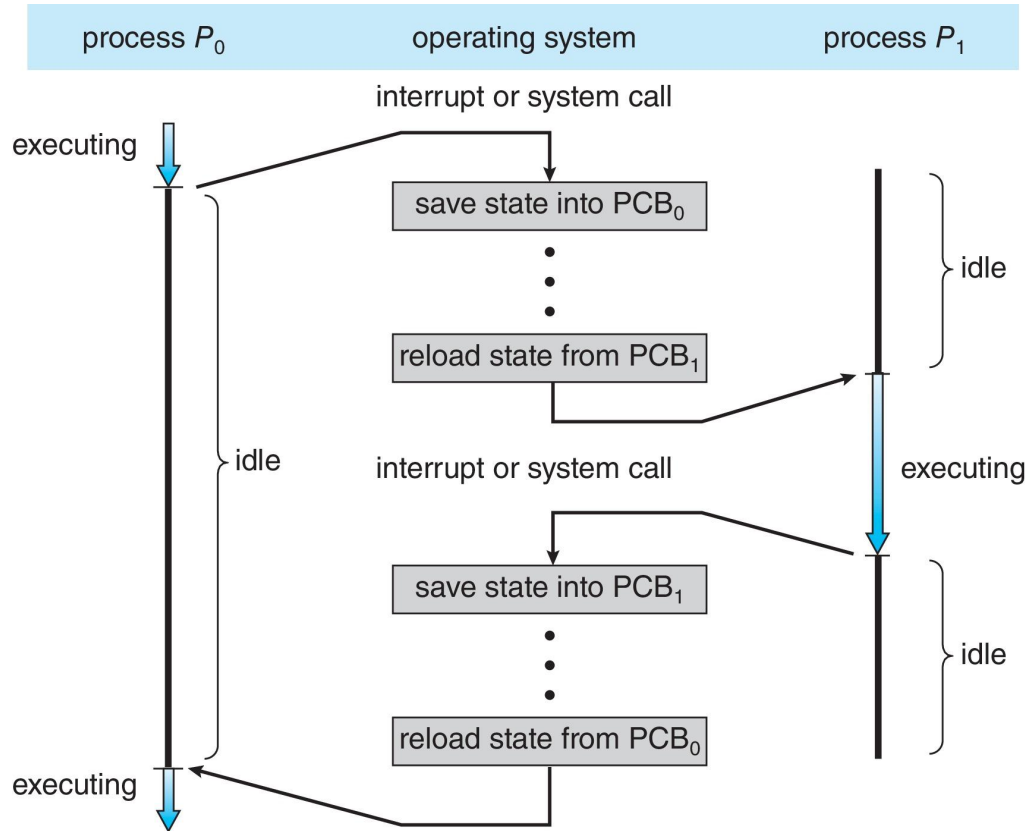
- Processes that are waiting for a certain event to occur — such as completion of I/O — are placed in a **wait queue**

# Representation of Process Scheduling



# CPU Switch From Process to Process

A **context switch** occurs when the CPU switches from one process to another.



# Context Switch

- When CPU switches to another process, the system must **save the state** of the old process and load the **saved state** for the new process via a **context switch**
- **Context** of a process represented in the PCB
- Context-switch time is pure overhead; the system does no useful work while switching
  - The more complex the OS and the PCB □ the longer the context switch
- Time dependent on hardware support
  - Some hardware provides multiple sets of registers per CPU □ multiple contexts loaded at once

# Multitasking in Mobile Systems

- Some mobile systems (e.g., early version of iOS) allow only one process to run, others suspended
- Due to screen real estate, user interface limits iOS provides for a
  - Single **foreground** process- controlled via user interface
  - Multiple **background** processes— in memory, running, but not on the display, and with limits
    - Limits include single, short task, receiving notification of events, specific long-running tasks like audio playback

A book on macOS

<https://newosxbook.com/home.html>

<https://newosxbook.com/jbooks.html>

<https://newosxbook.com/MOXil.pdf>



# Android details

- Android runs foreground and background, with fewer limits ([Processes and app lifecycle | Android Developers](#) )
- When deciding how to classify a process, the system bases its decision on the most important level found among all the components currently active in the process (see [Activity](#), [Service](#), and [BroadcastReceiver](#) ).
- 
- [Activity](#), [Service](#), and [BroadcastReceiver](#) impact the lifetime of the application's process.
  - system determines which processes to kill when low on memory

Based on [Activity](#), [Service](#), and [BroadcastReceiver](#)

- **Foreground, visible, service, or cached process**

- **A foreground process** is one that is required for what the user is currently doing.
  - It is running an [Activity](#) at the top of the screen that the user is interacting with
  - It has a [BroadcastReceiver](#) that is currently running
  - It has a [Service](#) that is currently executing code in one of its callbacks.
- **A visible process** is doing work that the user is currently aware of, so killing it has a noticeable negative impact on the user experience.
  - It is running an [Activity](#) that is visible to the user on-screen but not in the foreground
  - It has a [Service](#) that is running as a foreground service, through [Service.startForeground\(\)](#)
  - It is hosting a service that the system is using for a particular feature that the user is aware of, such as a live wallpaper or an input method service

- A **service process** is one holding a Service that has been started with the startService() method.
  - Though these processes are not directly visible to the user, they are generally doing things that the user cares about (such as background network data upload or download),
  - so the system always keeps such processes running unless there is not enough memory to retain all foreground and visible processes.
- A **cached process** is one that is not currently needed, so the system is free to kill it as needed when resources like memory are needed elsewhere.
  - In a normally behaving system, these are the only processes involved in resource management


# Operations on Processes

- System must provide mechanisms for:
  - Process creation
  - Process termination

# Process Creation


- **Parent** process create **children** processes, which, in turn create other processes, forming a **tree** of processes
- Generally, process identified and managed via a **process identifier (pid)**
- Resource sharing options
  - Parent and children share all resources
  - Children share subset of parent's resources
  - Parent and child share no resources
- Execution options
  - Parent and children execute concurrently
  - Parent waits until children terminate

# A Tree of Processes in Linux



# Process Creation (Cont.)

- Address space
  - Child duplicate of parent
  - Child has a program loaded into it
- UNIX examples
  - **fork()** system call creates new process
  - **exec()** system call used after a **fork()** to replace the process' memory space with a new program
  - Parent process calls **wait()** waiting for the child to terminate







# Program Forking Separate Process

```
#include < sys/types.h >
#include < stdio.h >
#include < unistd.h >

int main(){
    pid_t pid;
    /* fork a child process */
    pid = fork();
    if (pid < 0){ /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0){ /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else{ /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }
    return 0;
}
```

# Creating a Separate Process via Windows API

```
#include < stdio.h >
#include < windows.h >
int main(VOID) {
    STARTUPINFO si;
    PROCESS_INFORMATION pi;
    /* allocate memory */
    ZeroMemory(&si, sizeof(si));
    si.cb = sizeof(si);
    ZeroMemory(&pi, sizeof(pi));
    /* create child process */
    if (!CreateProcess(NULL, /* use command line */
        "C: \\ WINDOWS \\ system32 \\ mspaint.exe", /* command */
        NULL, /* don't inherit process handle */
        NULL, /* don't inherit thread handle */
        FALSE, /* disable handle inheritance */
        0, /* no creation flags */
        NULL, /* use parent's environment block */
        NULL, /* use parent's existing directory */
        &si, &pi)) {
        fprintf(stderr, "Create Process Failed");
        return -1;
    }
    /* parent will wait for the child to complete */
    WaitForSingleObject(pi.hProcess, INFINITE);
    printf("Child Complete");

    /* close handles */
    CloseHandle(pi.hProcess);
    CloseHandle(pi.hThread);
}
```

# Process Termination

- Process executes last statement and then asks the operating system to delete it using the `exit()` system call.
  - Returns status data from child to parent (via `wait()`)
  - Process' resources are deallocated by operating system
- Parent may terminate the execution of children processes using the `kill()` system call. Some reasons for doing so:
  - Child has exceeded allocated resources
  - Task assigned to child is no longer required
  - The parent is exiting, and the operating system does not allow a child to continue if its parent terminates

●

# Process Termination

- Some operating systems do not allow child to exist if its parent has terminated.
- If a process terminates, then all its children must also be terminated.
  - **cascading termination.** All children, grandchildren, etc., are terminated.
  - The termination is initiated by the operating system.
- The parent process may wait for termination of a child process by using the `wait()` system call.
- The call returns status information and the pid of the terminated process

```
pid = wait(&status);
```

- If no parent waiting (did not invoke `wait()`) process is a **zombie**
- If parent terminated without invoking `wait()`, process is an **orphan**

# Android Process Importance Hierarchy

- Mobile operating systems often have to terminate processes to reclaim system resources such as memory. From **most** to **least** important:
  - Foreground process
  - Visible process
  - Service process
  - Background process
  - Empty process
- Android will begin terminating processes that are least important.
  - **the states are saved before the termination**

# Interprocess Communication (IPC)

# Multiprocess Architecture – Chrome Browser



- Many web browsers ran as single process (some still do)
  - If one web site causes trouble, entire browser can hang or crash
- Google Chrome Browser is multiprocess with 3 different types of processes:
  - **Browser** process manages user interface, disk and network I/O
  - **Renderer** process renders web pages, deals with HTML, Javascript.
    - A new renderer created for each website opened
    - Runs in **sandbox** restricting disk and network I/O, minimizing effect of security exploits
  - **Plug-in** process for each type of plug-in

<https://www.chromium.org/developers/design-documents/multi-process-architecture/>


# Interprocess Communication

- Processes within a system may be *independent* or *cooperating*
- Cooperating process can affect or be affected by other processes, including sharing data
- Reasons for cooperating processes:
  - Information sharing
  - Computation speedup
  - Modularity
  - Convenience
- Cooperating processes need **interprocess communication (IPC)**



# Two models of IPC

(a) **Shared memory.**

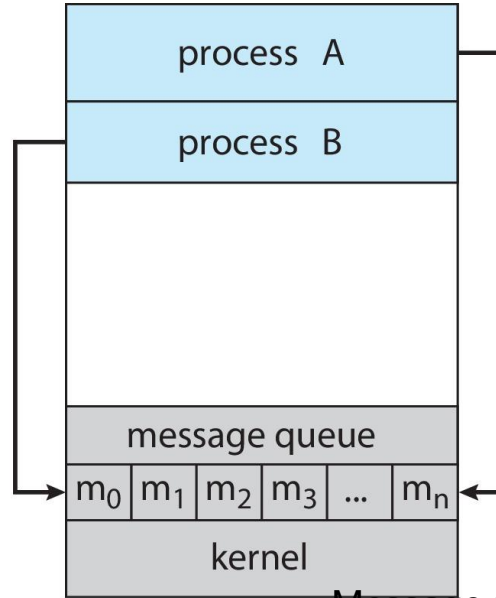


## Shared memory

(a)

- fast, efficient(no overhead), ideal for large data sharing
- needs synchronization, security, management

(b) **Message passing.**



(b)

## Message passing

- pros: secure, flexible(remote or local), error handling
- cons: Latency, overhead, complexity

# Producer-Consumer Problem

- Paradigm for cooperating processes:
  - *producer* process produces information that is consumed by a *consumer* process
- Two variations:
  - **unbounded-buffer** places no practical limit on the size of the buffer:
    - Producer never waits
    - Consumer waits if there is no buffer to consume
  - **bounded-buffer** assumes that there is a fixed buffer size
    - Producer must wait if all buffers are full
    - Consumer waits if there is no buffer to consume

## IPC – Shared Memory

- An area of memory shared among the processes that wish to communicate
- The communication is under the control of the users processes not the operating system.
- Major issues is to provide mechanism that will allow the user processes to synchronize their actions when they access shared memory.

# Bounded-Buffer – Shared-Memory Solution

- Shared data

```
#define BUFFER_SIZE 10
typedef struct {
    . . .
} item;

item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
```

- Solution is correct, but can only use `BUFFER_SIZE-1` elements

# Producer Process – Shared Memory

```
item next_produced;

while (true) {
    /* produce an item in next produced */
    while (((in + 1) % BUFFER_SIZE) == out)
        ; /* do nothing */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
}
```

# Consumer Process – Shared Memory

```
item next_consumed;  
while (true) {  
    while (in == out)  
        ; /* do nothing */  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
  
    /* consume the item in next consumed */  
}
```

- this illustration does not address concerns the situation in which both the producer process and the consumer process attempt to access the shared buffer concurrently.

# What about Filling all the Buffers?

- consumer-producer problem fills **all** the buffers.
  - an integer **counter** that keeps track of the number of full buffers.
  - Initially, **counter** is set to 0.
  - The integer **counter** is incremented by the producer after it produces a new buffer.
  - The integer **counter** is and is decremented by the consumer after it consumes a buffer.

# Producer

```
while (true) {  
    /* produce an item in next  
    produced */  
  
    while (counter == BUFFER_SIZE)  
        ; /* do nothing */  
  
    buffer[in] = next_produced;  
    in = (in + 1) % BUFFER_SIZE;  
    counter++;  
  
}
```

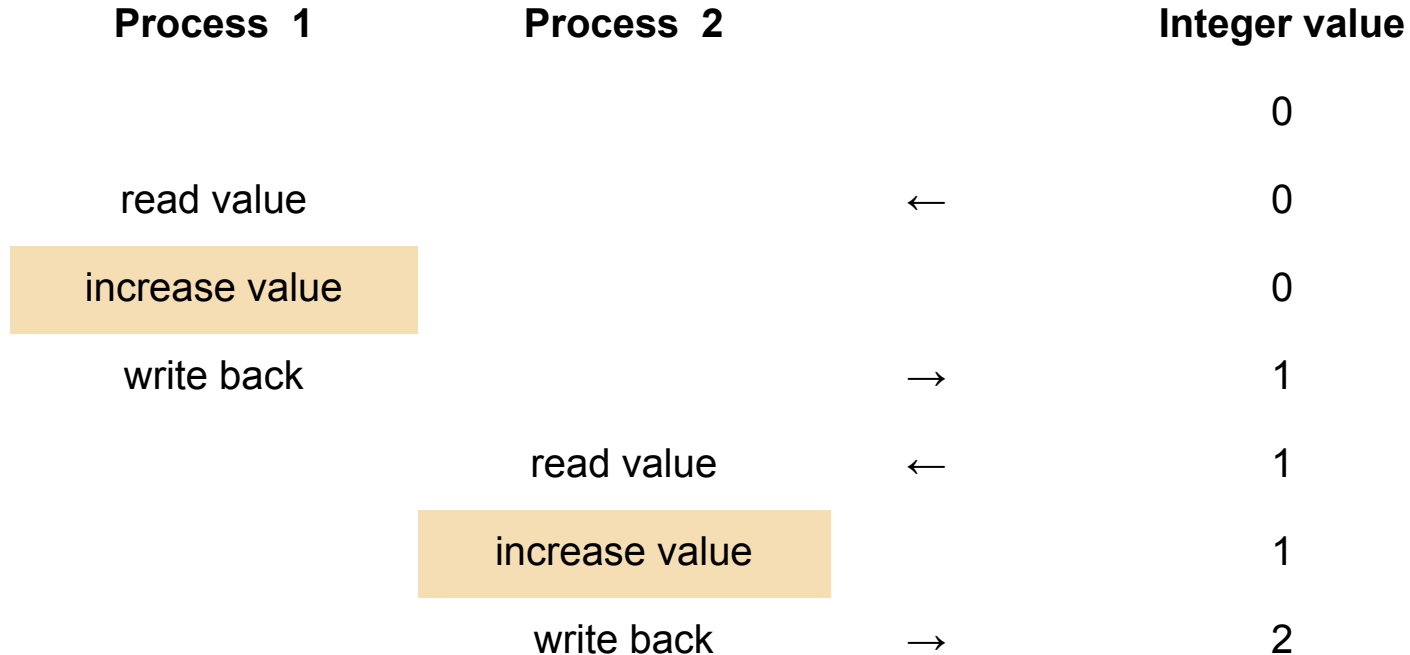
# Consumer

```
while (true) {  
    while (counter == 0)  
        ; /* do nothing */  
  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
  
    counter--;  
  
    /* consume the item in next  
    consumed */  
  
}
```



# Race Condition

[Race condition - Wikipedia](#) is the condition of an electronics, software, or other system where the system's substantive behavior is dependent on the sequence or timing of other uncontrollable events.



# Race Condition

**counter++** could be implemented as

```
register1 = counter
register1 = register1 + 1
counter = register1
```

**counter--** could be implemented as

```
register2 = counter
register2 = register2 - 1
counter = register2
```

- Consider this execution interleaving with “count = 5” initially:


S0: producer execute	<code>register1 = counter</code>	{register1 = 5}
S1: producer execute	<code>register1 = register1 + 1</code>	{register1 = 6}
S2: consumer execute	<code>register2 = counter</code>	{register2 = 5}
S3: consumer execute	<code>register2 = register2 - 1</code>	{register2 = 4}
S4: producer execute	<code>counter = register1</code>	{counter = 6}
S5: consumer execute	<code>counter = register2</code>	{counter = 4}

# Race Condition (Cont.)

- Question - ~~why was there no race condition in the first solution (where at most  $N - 1$  buffers can be filled?)~~
- Her iki çözümde de race condition var. Bu kısım tam doğru değil!
  - e.g., 1 den fazla producer ve/veya 1den fazla consumer
  - yine read/write arasında sync olmadığı için, sıraları değişebilir



# IPC – Message Passing



- Processes communicate with each other without resorting to shared variables
- IPC facility provides two operations:
  - `send(destination, &message);`
  - `receive(source, &message);`
- The *message* size is either fixed or variable

# Message Passing (Cont.)

- If processes  $P$  and  $Q$  wish to communicate, they need to:
  - Establish a **communication link** between them
  - Exchange messages via send/receive
- Implementation issues:
  - How are links established?
  - Can a link be associated with more than two processes?
  - How many links can there be between every pair of communicating processes?
  - What is the capacity of a link?
  - Is the size of a message that the link can accommodate fixed or variable?
  - Is a link unidirectional or bi-directional?

# Implementation of Communication Link

- Physical:
  - Shared memory
  - Hardware bus
  - Network
- Logical:
  - Direct or indirect
  - Synchronous or asynchronous
  - Automatic or explicit buffering

# Direct Communication

- Processes must name each other explicitly:
  - `send(P, message)` – send a message to process P
  - `receive(Q, message)` – receive a message from process Q
- Properties of communication link
  - Links are established automatically
  - A link is associated with exactly one pair of communicating processes
  - Between each pair there exists exactly one link
  - The link may be unidirectional, but is usually bi-directional

# Indirect Communication

- Messages are directed and received from mailboxes (also referred to as ports)
  - Each mailbox has a unique id
  - Processes can communicate only if they share a mailbox
- Properties of communication link
  - Link established only if processes share a common mailbox
  - A link may be associated with many processes
  - Each pair of processes may share several communication links
  - Link may be unidirectional or bi-directional



# Indirect Communication (Cont.)

- Operations
  - Create a new mailbox (port)
  - Send and receive messages through mailbox
  - Delete a mailbox
- Primitives are defined as:
  - **send**(*A, message*) – send a message to mailbox A
  - **receive**(*A, message*) – receive a message from mailbox A

# Indirect Communication (Cont.)

- Mailbox sharing

- $P_1$ ,  $P_2$ , and  $P_3$  share mailbox A
- $P_1$  sends;  $P_2$  and  $P_3$  receive
- Who gets the message?

- Solutions

- Allow a link to be associated with at most two processes
- Allow only one process at a time to execute a receive operation
- Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was.

# Synchronization

Message passing may be either blocking or non-blocking

- **Blocking** is considered **synchronous**
  - **Blocking send** -- the sender is blocked until the message is received
  - **Blocking receive** -- the receiver is blocked until a message is available
- **Non-blocking** is considered **asynchronous**
  - **Non-blocking send** -- the sender sends the message and continue
  - **Non-blocking receive** -- the receiver receives:
    - A valid message, or
    - Null message
- Different combinations possible
  - If both send and receive are blocking, we have a **rendezvous**

# Producer-Consumer: Message Passing

- Producer

```
message next_produced;  
while (true) {  
    /* produce an item in next_produced */  
  
    send(next_produced);  
}
```

- Consumer

```
message next_consumed;  
while (true) {  
    receive(next_consumed)  
  
    /* consume the item in next_consumed */  
}
```

# Buffering

- Queue of messages attached to the link.
- Implemented in one of three ways
  1. Zero capacity – no messages are queued on a link.  
Sender must wait for receiver (rendezvous)
  2. Bounded capacity – finite length of  $n$  messages  
Sender must wait if link full
  3. Unbounded capacity – infinite length  
Sender never waits

# Examples of IPC Systems - POSIX

- POSIX Shared Memory
  - Process first creates shared memory segment  
`shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);`
  - Also used to open an existing segment
  - Set the size of the object  
`ftruncate(shm_fd, 4096);`
  - Use `mmap()` to memory-map a file pointer to the shared memory object
  - Reading and writing to shared memory is done by using the pointer returned by `mmap()`.

# IPC POSIX Producer

```
#include <stdio.h>#include<stdlib.h>#include<string.h>#include<fcntl.h>
#include <sys/shm.h>#include<sys/stat.h>#include <sys/mman.h>
int main() {
    const int SIZE = 4096; /* the size (in bytes) of shared memory object */
    const char *name = "OS"; /* name of the shared memory object */

    /* strings written to shared memory */
    const char *message0 = "Hello", *message1 = "World!";
    int fd; /* shared memory file descriptor */

    char *ptr; /* pointer to shared memory object */

    /* create the shared memory object */
    fd = shm_open(name, O_CREAT | O_RDWR, 0666);
    /* configure the size of the shared memory object */
    ftruncate(fd, SIZE);
    /* memory map the shared memory object */
    ptr = (char *)mmap(0, SIZE, PROT_READ | PROT_WRITE, MAP_SHARED, fd, 0);

    /* write to the shared memory object */
    sprintf(ptr, "%s", message0);
    ptr += strlen(message0);
    sprintf(ptr, "%s", message1);
    ptr += strlen(message1);
    return 0;
```

```
}
```

# IPC POSIX Consumer

```
#include < sys/mman.h>
#include <fcntl.h>
#include <stdio.h>
#include <stdlib.h>
#include <sys/shm.h>
#include <sys/stat.h>
int main() {
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* shared memory file descriptor */
    int fd;
    /* pointer to shared memory object */
    char *ptr;
    /* open the shared memory object */
    fd = shm_open(name, O_RDONLY, 0666);
    /* memory map the shared memory object */
    ptr = (char *)mmap(0, SIZE, PROT_READ | PROT_WRITE, MAP_SHARED, fd, 0);
    /* read from the shared memory object */
    printf("%s", (char *)ptr);
    /* remove the shared memory object */
    shm_unlink(name);
    return 0;
}
```



# Examples of IPC Systems - Mach

[Mach Project Publications and Related Documents](#)

[The GNU Mach Reference Manual](#)

- Mach communication is message based
  - Even system calls are messages
  - Each task gets two ports at creation - Kernel and Notify
  - Messages are sent and received using the `mach_msg()` function
- Ports needed for communication, created via  
`mach_port_allocate()`
- Send and receive are flexible; for example four options if mailbox full:
  - Wait indefinitely
  - Wait at most n milliseconds
  - Return immediately
  - Temporarily cache a message

# Mach Messages

```
#include<mach/mach.h>

struct message {
    mach_msg_header_t header;
    int data;
};

mach_port_t client;
mach_port_t server;
```

# Mach Message Passing - Client

```
mach_msg_return_t mach_msg(  
    mach_msg_header_t *msg, mach_msg_option_t option,  
    mach_msg_size_t send_size, mach_msg_size_t rcv_size,  
    mach_port_t rcv_name, mach_msg_timeout_t timeout,  
    mach_port_t notify)  
  
/* Client Code */  
  
struct message message;  
  
// construct the header  
message.header.msgh_size = sizeof(message);  
message.header.msgh_remote_port = server;  
message.header.msgh_local_port = client;  
  
// send the message  
mach_msg(&message.header, // message header  
    MACH_SEND_MSG, // sending a message  
    sizeof(message), // size of message sent  
    0, // maximum size of received message - unnecessary  
    MACH_PORT_NULL, // name of receive port - unnecessary  
    MACH_MSG_TIMEOUT_NONE, // no time outs  
    MACH_PORT_NULL // no notify port  
);
```

# Mach Message Passing - Server



```
/* Server Code */

struct message message;

// receive the message
mach_msg(&message.header, // message header
        MACH_RCV_MSG, // sending a message
        0, // size of message sent
        sizeof(message), // maximum size of received message
        server, // name of receive port
        MACH_MSG_TIMEOUT_NONE, // no time outs
        MACH_PORT_NULL // no notify port
);
```

# Mach another example

```
kern_return_t err;  
mach_port_t rcv_port;  
  
err = mach_port_allocate(mach_task_self(),  
                        MACH_PORT_RIGHT_RECEIVE,  
                        &rcv_port);  
  
if (err != KERN_SUCCESS) {  
    perror("error : could not allocate any port\n");  
    exit(err);  
}  
  
struct integer_message {  
    mach_msg_header_t head;  
    mach_msg_type_t type;  
    int inline_integer;  
};
```



```

void
send_integer( mach_port_t destination, int i ){
    kern_return_t err;
    struct integer_message message;
    /* (i.a) Fill the header fields : */
    message.head.msgh_bits =
        MACH_MSGH_BITS_REMOTE(MACH_MSG_TYPE_MAKE_SEND);
    message.head.msgh_size = sizeof( struct integer_message );
    message.head.msgh_local_port = MACH_PORT_NULL;
    message.head.msgh_remote_port = destination;

    struct integer_message{
        mach_msg_header_t head;
        mach_msg_type_t type;
        int inline_integer;
    };

    /* (i.b) Explain the message type ( an integer ) */
    message.type.msgt_name = MACH_MSG_TYPE_INTEGER_32;
    message.type.msgt_size = 32;
    message.type.msgt_number = 1;
    message.type.msgt_inline = TRUE;
    message.type.msgt_longform = FALSE;
    message.type.msgt_deallocate = FALSE;
    /* message.type.msgt_unused = 0; */ /* not needed, I think */
    /* (i.c) Fill the message with the given integer : */
    message.inline_integer = i;

```

```

/* (ii) Send the message : */

```

```

void
send_integer( mach_port_t destination, int i ){
    ...
    /* (ii) Send the message : */
    err = mach_msg( &(message.head), MACH_SEND_MSG,
                    message.head.msgh_size, 0, MACH_PORT_NULL,
                    MACH_MSG_TIMEOUT_NONE, MACH_PORT_NULL );
    /* (iii) Analysis of the error code;
    if succes, print and acknowledge message and return */
    if( err == MACH_MSG_SUCCESS ){
        printf( "success: the message was queued\n" );
    }
    else{
        perror( "error: some unexpected error ocurred!\n" );
        exit(err);
    }

    return;
}

```

[https://hurdextras.nongnu.org/ipc\\_guide/mach\\_ipc\\_basic\\_concepts.html](https://hurdextras.nongnu.org/ipc_guide/mach_ipc_basic_concepts.html)

# Receive integer

```
void receive_integer(mach_port_t source, int *ip) {
    kern_return_t err;

    struct integer_message message;

    /* (i) Fill the little thing we know about the message : */
    /* message.head.msgh_size = sizeof(struct integer_message ); */

    /* (ii) Receive the message : */
    err = mach_msg(&(message.head), MACH_RCV_MSG, 0, message.head.msgh_size,
                   source, MACH_MSG_TIMEOUT_NONE, MACH_PORT_NULL);

    if (err == MACH_MSG_SUCCESS) {
        printf("success: the message was received\n");
    } else {
        perror("error: Some unexpected error occurred\n");
        exit(err);
    }

    *ip = message.inline_integer;

    return;
}
```


[https://hurdextras.nongnu.org/ipc\\_guide/mach\\_ipc\\_basic\\_concepts.html](https://hurdextras.nongnu.org/ipc_guide/mach_ipc_basic_concepts.html)



# Examples of IPC Systems – Windows

- Message-passing centric via **advanced local procedure call (LPC)** facility
  - Only works between processes on the same system
  - Uses ports (like mailboxes) to establish and maintain communication channels
  - Communication works as follows:
    - Server(subsystem) processes publish **connection port** objects that are visible to all processes.
    - The client requests a connection to named port.
    - The server creates two private **communication ports**
      - **client communication port**
        - returns the handle to the client.
      - **server communication port**
    - The client and server use the corresponding port handle to send messages or callbacks and to listen for replies.

# Local Procedure Calls in Windows




# Pipes

- allows two or more processes to communicate with each other by creating a unidirectional or bidirectional channel between them
- Issues:
  - Is communication unidirectional or bidirectional?
  - In the case of two-way communication, is it half or full-duplex?
  - Must there exist a relationship (i.e., **parent-child**) between the communicating processes?
  - Can the pipes be used over a network?
- **Ordinary pipes** – cannot be accessed from outside the process that created it. Typically, a parent process creates a pipe and uses it to communicate with a child process that it created.
- **Named pipes** – can be accessed without a parent-child relationship.

# Ordinary Pipes

- Ordinary Pipes allow communication in standard producer-consumer style
- Producer writes to one end (the **write-end** of the pipe)
- Consumer reads from the other end (the **read-end** of the pipe)
- Ordinary pipes are therefore unidirectional
- Require parent-child relationship between communicating processes



- Windows calls these **anonymous pipes**

# Named Pipes

- Named Pipes are more powerful than ordinary pipes
- Communication is bidirectional
- No parent-child relationship is necessary between the communicating processes
- Several processes can use the named pipe for communication
- Provided on both UNIX and Windows systems


# Communications in Client-Server Systems

- Sockets
- Remote Procedure Calls

# Sockets

- A **socket** is defined as an endpoint for communication
- Concatenation of IP address and **port** – a number included at start of message packet to differentiate network services on a host
- The socket **161.25.19.8:1625** refers to port **1625** on host **161.25.19.8**
- Communication consists between a pair of sockets
- All ports below 1024 are ***well known***, used for standard services
- Special IP address 127.0.0.1 (**loopback**) to refer to system on which process is running

# Socket Communication





# Sockets in Java

- Three types of sockets
  - **Connection-oriented (TCP)**
  - **Connectionless (UDP)**
  - `MulticastSocket` class— data
- Consider this “Date” server in

```
import java.net.*;
import java.io.*;

public class DateServer
{
    public static void main(String[] args) {
        try {
            ServerSocket sock = new ServerSocket(6013);

            /* now listen for connections */
            while (true) {
                Socket client = sock.accept();

                PrintWriter pout = new
                    PrintWriter(client.getOutputStream(), true);

                /* write the Date to the socket */
                pout.println(new java.util.Date().toString());

                /* close the socket and resume */
                /* listening for connections */
                client.close();
            }
        }
        catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
```

# Sockets in Java

## The equivalent Date client

```
import java.net.*;
import java.io.*;

public class DateClient
{
    public static void main(String[] args) {
        try {
            /* make connection to server socket */
            Socket sock = new Socket("127.0.0.1",6013);

            InputStream in = sock.getInputStream();
            BufferedReader bin = new
                BufferedReader(new InputStreamReader(in));

            /* read the date from the socket */
            String line;
            while ( (line = bin.readLine()) != null)
                System.out.println(line);

            /* close the socket connection*/
            sock.close();
        }
        catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
```

# Remote Procedure Calls

- Remote procedure call (RPC) abstracts procedure calls between processes on networked systems
  - Again uses ports for service differentiation
- **Stubs** – client-side proxy for the actual procedure on the server
- The client-side stub locates the server and **marshalls** the parameters
- The server-side stub receives this message, unpacks the marshalled parameters, and performs the procedure on the server
- On Windows, stub code compile from specification written in **Microsoft Interface Definition Language (MIDL)**

# Remote Procedure Calls (Cont.)

- Data representation handled via **External Data Representation (XDL)** format to account for different architectures
  - **Big-endian** and **little-endian**
- Remote communication has more failure scenarios than local
  - Messages can be delivered ***exactly once*** rather than ***at most once***
- OS typically provides a rendezvous (or **matchmaker**) service to connect client and server

